

Exploratory Data Analysis

Matheus C. Andrade

6/15/2021

Stage 0: Collecting the data

```
## Setting the working directory
setwd("~/Development/DataScienceAcademy/FCD/BigDataRAzure/ProjetoFinal/TalkingData-AdTracking-Fraud-Detection")
getwd()

## [1] "/home/matheus/Development/DataScienceAcademy/FCD/BigDataRAzure/ProjetoFinal/TalkingData-AdTracking-Fraud-Detection"

## Libraries
library(data.table)
library(Amelia)

## Loading required package: Rcpp
## ##
## ## Amelia II: Multiple Imputation
## ## (Version 1.8.0, built: 2021-05-26)
## ## Copyright (C) 2005-2021 James Honaker, Gary King and Matthew Blackwell
## ## Refer to http://gking.harvard.edu/amelia/ for more information
## ##

library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:data.table':
##
##   between, first, last

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(ggplot2)
library(gridExtra)

##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
```

```
##      combine
library(corrplot)

## corrplot 0.88 loaded
## Loading dataset
df_original <- fread("../Datasets/train_sample.csv", header=T)
# df <- fread(file.choose(), header=T)
df <- df_original
```

Stage 1: Knowing the data as they are

```
str(df) #glimpse(df)

## Classes 'data.table' and 'data.frame':  100000 obs. of  8 variables:
## $ ip      : int  87540 105560 101424 94584 68413 93663 17059 121505 192967 143636 ...
## $ app     : int  12 25 12 13 12 3 1 9 2 3 ...
## $ device  : int  1 1 1 1 1 1 1 1 2 1 ...
## $ os      : int  13 17 19 13 1 17 17 25 22 19 ...
## $ channel : int  497 259 212 477 178 115 135 442 364 135 ...
## $ click_time : POSIXct, format: "2017-11-07 09:30:38" "2017-11-07 13:40:27" ...
## $ attributed_time: POSIXct, format: NA NA ...
## $ is_attributed : int  0 0 0 0 0 0 0 0 0 0 ...
## - attr(*, ".internal.selfref")=<externalptr>

dim(df)

## [1] 100000      8

summary(df)

##           ip           app           device           os
## Min.      :    9   Min.      : 1.00   Min.      : 0.00   Min.      : 0.00
## 1st Qu.: 40552   1st Qu.:  3.00   1st Qu.:  1.00   1st Qu.: 13.00
## Median : 79827   Median : 12.00   Median :  1.00   Median : 18.00
## Mean      : 91256   Mean      : 12.05   Mean      : 21.77   Mean      : 22.82
## 3rd Qu.:118252   3rd Qu.: 15.00   3rd Qu.:  1.00   3rd Qu.: 19.00
## Max.      :364757   Max.      :551.00   Max.      :3867.00   Max.      :866.00
##
##           channel      click_time      attributed_time
## Min.      :  3.0   Min.      :2017-11-06 16:00:00   Min.      :2017-11-06 17:19:04
## 1st Qu.:145.0   1st Qu.:2017-11-07 11:34:09   1st Qu.:2017-11-07 11:50:27
## Median :258.0   Median :2017-11-08 07:07:50   Median :2017-11-08 06:43:39
## Mean      :268.8   Mean      :2017-11-08 06:29:52   Mean      :2017-11-08 07:04:12
## 3rd Qu.:379.0   3rd Qu.:2017-11-09 02:06:01   3rd Qu.:2017-11-09 01:42:52
## Max.      :498.0   Max.      :2017-11-09 15:59:51   Max.      :2017-11-09 15:28:15
##
##                                     NA's      :99773
## is_attributed
## Min.      :0.00000
## 1st Qu.:0.00000
## Median :0.00000
## Mean      :0.00227
## 3rd Qu.:0.00000
## Max.      :1.00000
##
```

```
## Duplicated rows analysis
any(duplicated(df))

## [1] TRUE

# what are this rows?
df[duplicated(df), ] #df %>% !distinct() # There is one

##      ip app device os channel      click_time attributed_time is_attributed
## 1: 871  12      1 13      178 2017-11-08 10:00:05      <NA>              0

# removing duplicated rows
df <- df[!duplicated(df), ]
any(duplicated(df))

## [1] FALSE

## quantity of null values in each columns
any(is.na(df))

## [1] TRUE

# missmap(df, main = "Missing Values Map", col = c("red", "black"), legend = FALSE)
sapply(df, function(x) sum(is.na(x)))

##           ip           app           device           os           channel
##           0           0           0           0           0
## click_time attributed_time is_attributed
##           0          99772           0

# just 'attributed_time' has null values: 99773
sum(is.na(df[, "attributed_time"]))/length(df$attributed_time)*100

## [1] 99.773

# as for the users who didn't download the app, the time has not been recorded
# and the column wasn't filled with any value

## quantity of unique values in each columns
sapply(df, function(x) length(unique(x)))

##           ip           app           device           os           channel
##          34857          161           100          130           161
## click_time attributed_time is_attributed
##          80350          228           2

# the label column is a factor with two levels
df$is_attributed <- as.factor(df$is_attributed)

# "ip", "app", "device", "os", "channel" are also factor type
df[,1:5] <- lapply(df[,1:5], factor)#####
#str(df)

## High class imbalance problem
summary(df$is_attributed) #table(df$is_attributed)

##      0      1
## 99772  227
```

```

prop.table(table(df$is_attributed))*100

##
##           0           1
## 99.7729977  0.2270023
# it means that the models will be overfitted about no-downloads

## when "attributed_time" is NULL "is_attributed" is 0.
# there isn't attributed_time when there wasn't made download

## Generating weekday and hour from click_time intending to explore days and hours
# 'weekdays' to names, 'wday' to numbers (it starts with 0 = Sunday)
df$weekday <- weekdays(df$click_time)
df$hour <- hour(df$click_time)
#glimpse(df)
# quantity of unique values in the new columns
sapply(df[, 9:10], function(x) length(unique(x)))

## weekday      hour
##          4       24
## changing the columns order
# names(df)
df <- df[, c(6,9,10,1,2,3,4,5,7,8)]

# changing for factor the new columns
df$weekday <- as.factor(df$weekday)
df$hour <- as.factor(df$hour)

# after data munging, just confirming whether the data integrity is as before
sapply(df, function(x) sum(is.na(x)))

##      click_time      weekday      hour      ip      app
##           0           0           0           0           0
##      device      os      channel attributed_time is_attributed
##           0           0           0          99772           0
sapply(df, function(x) length(unique(x)))

##      click_time      weekday      hour      ip      app
##      80350           4           24      34857      161
##      device      os      channel attributed_time is_attributed
##      100          130          161          228           2

## Creating subsets from is_attributed classes
df_IsAttributed0 <- df %>%
  filter(is_attributed == '0')

df_IsAttributed1 <- df %>%
  filter(is_attributed == '1') %>%
  mutate(wday_IsAttributed1 = weekdays(attributed_time),
         hour_IsAttributed1 = hour(attributed_time))

# as attributed_time represents the event time we want predict, it's deleted
df$attributed_time = NULL

```

```
summary(df)
```

```
##      click_time      weekday      hour
## Min.   :2017-11-06 16:00:00 Monday   : 5011  4      : 6039
## 1st Qu.:2017-11-07 11:34:09 Thursday :28561 0      : 5654
## Median :2017-11-08 07:07:49 Tuesday   :32393 13     : 5619
## Mean   :2017-11-08 06:29:52 Wednesday:34034 14     : 5561
## 3rd Qu.:2017-11-09 02:06:01           10     : 5510
## Max.   :2017-11-09 15:59:51           5      : 5400
##                                     (Other):66216
##      ip      app      device      os
## 5348 : 669  3      :18279  1      :94337  19     :23870
## 5314 : 616  12     :13197  2      : 4345  13     :21222
## 73487 : 439  2      :11737  0      :  541  17     : 5232
## 73516 : 399  9      : 8992 3032   :  371  18     : 4830
## 53454 : 280  15     : 8595 3543   :  151  22     : 4039
## 114276 : 219  18     : 8315 3866   :  93   10     : 2816
## (Other):97377 (Other):30884 (Other): 161 (Other):37990
##      channel      is_attributed
## 280      : 8114  0:99772
## 245      : 4802  1: 227
## 107      : 4543
## 477      : 3960
## 134      : 3224
## 259      : 3130
## (Other):72226
```

Stage 2: Exploratory Data Analysis

2.1 Bar plots

```
# For categorical variables (or grouping variables).
# the count of categories is visualized using a bar plot, pie chart or dot charts to show the proportion

# creating a list with the same dataset rows length
df_subsets <- list(1:nrow(df))
df_subsets[[1]][1]

# creting a aux dataset with only columns in df_subsets
df_aux <- df[, 2:8] # without click_time, which is represented by hour and weekday columns

### 2.1.1 Creating dataset for each variable counting its frequency
for(i in 1:7){
  #print(names(df_dataSets[i]) )
  df_subsets[[i]] <- df_aux %>%
    group_by_at(i) %>%
    summarise(counts = n())
}
#View(df_subsets)

### 2.1.2 Creating bar plots for each variable counting its frequency from its datasets
```

```

# automatizing visualization for bar plots
for (i in 1:length(df_subsets)){


  if(i==3)
    next # as the IP bar plot is heavy for load, this bar is omitted


  x_column <- unlist(df_subsets[[i]][,1])
  title_x <- names(df_subsets[[i]][,1])
  title_y <- names(df_subsets[[i]][,2])

  barplot <- ggplot(df_subsets[[i]],aes(x = x_column , y = counts)) +
    geom_bar(fill = "#00A4DEF7", stat = "identity") +
    #geom_text(aes(label = counts), vjust = -0.3, size = 3) +
    ggtitle(paste("Bar Plot", i, title_x,"x", title_y)) +
    theme(plot.title = element_text(hjust = 0.5),
          #axis.title.x=element_blank(),
          axis.title.y=element_blank()) +
    xlab(title_x)


  print(barplot +
    if (i == 3)
      theme(axis.text.x = element_text(angle=-90, size=0))
    else if (i == 4)
      theme(axis.text.x = element_text(angle=-90, size=3))
    else if (i == 5)
      theme(axis.text.x = element_text(angle=-90, size=5))
    else if (i == 6)
      theme(axis.text.x = element_text(angle=-90, size=4))
    else if (i == 7)
      theme(axis.text.x = element_text(angle=-90, size=2.5)))
}

```







Bar Plot 4 app x counts




Bar Plot 5 device x counts



Bar Plot 6 os x counts



Bar Plot 7 channel x counts



Day: wednesday > tuesday > thursday > monday (a lot less)

Hour: between 17 and 22 hour is made less download; great peak: 0h to 15h; peak 1 is 0h to 7h; peak 2 is 9h to 15h; curiously 8 has the minor value compared with its proximities; @ that's interesting to relate day and hour together;

IP: there are 5 great used IP; there are also more +-20 less used than this 5 and more than others; @ to relate IP with (all) others variables?

App: in general, the firsts are more downloaded, but there is some whose highlight a tiny and one who highlight a lot more;

Device: there is one who highlights in a huge way;

OS: there are 2 who highlight a lot;

Channel: there is one that stands out a lot and others who highlights less;

2.2 Exploring hour and weekday

```
### 2.2.1 is_attributed frequency: comparison between download and no-download
# is_attributed0 weekday
df_IsAttributed0Day <- df_IsAttributed0 %>%
  group_by(weekday) %>%
  summarise(counts = n())

p1_IsA_Day <- ggplot(df_IsAttributed0Day, aes(x = weekday, y = counts)) +
  geom_bar(fill = "#00A4DEF7", stat = "identity") +
  geom_text(aes(label = counts), vjust = 1) +
```

```



ggtitle(" Hour when df_IsAttributed is 0") +
theme(plot.title = element_text(hjust = 0.5),
      axis.title.x=element_blank(),
      axis.title.y=element_blank())

# is_attributed1 weekday
df_IsAttributed1Day <- df_IsAttributed1 %>%
  group_by(weekday) %>%
  summarise(counts = n())

p2_IsA_Day <- ggplot(df_IsAttributed1Day, aes(x = weekday, y = counts)) +
  geom_bar(fill = "#00A4DEF7", stat = "identity") +
  geom_text(aes(label = counts), vjust = 1) +
  ggtitle(" Hour when df_IsAttributed is 1") +
  theme(plot.title = element_text(hjust = 0.5),
        axis.title.x=element_blank(),
        axis.title.y=element_blank())

grid.arrange(p1_IsA_Day, p2_IsA_Day, ncol = 1)

```



```

# Apparently there isn't no significant difference

# is_attributed0 hour
df_IsAttributed0Hour <- df_IsAttributed0 %>%
  group_by(hour) %>%
  summarise(counts = n())

```

```

p1_IsA_Hour <- ggplot(df_IsAttributed0Hour, aes(x = hour, y = counts)) +
  geom_bar(fill = "#00A4DEF7", stat = "identity") +
  geom_text(aes(label = counts), vjust = 1, size = 1.3) +
  ggtitle(" Hour when df_IsAttributed is 0") +
  theme(plot.title = element_text(hjust = 0.5),
        axis.title.x=element_blank(),
        axis.title.y=element_blank(),
        axis.text.x = element_text(size=5)) +
  geom_line(aes(x = hour, y = mean(counts)), size = 0.2, color="#8B4513", group = 1)


# is_attributed1 hour
df_IsAttributed1Hour <- df_IsAttributed1 %>%
  group_by(hour) %>%
  summarise(counts = n())

p2_IsA_Hour <- ggplot(df_IsAttributed1Hour, aes(x = hour, y = counts)) +
  geom_bar(fill = "#00A4DEF7", stat = "identity") +
  geom_text(aes(label = counts), vjust = 1, size = 3) +
  ggtitle(" Hour when df_IsAttributed is 1") +
  theme(plot.title = element_text(hjust = 0.5),
        axis.title.x=element_blank(),
        axis.title.y=element_blank(),
        axis.text.x = element_text(size=5)) +
  geom_line(aes(x = hour, y = mean(counts)), size = 0.2, color="#8B4513", group = 1)


grid.arrange(p1_IsA_Hour, p2_IsA_Hour, ncol = 1)

```

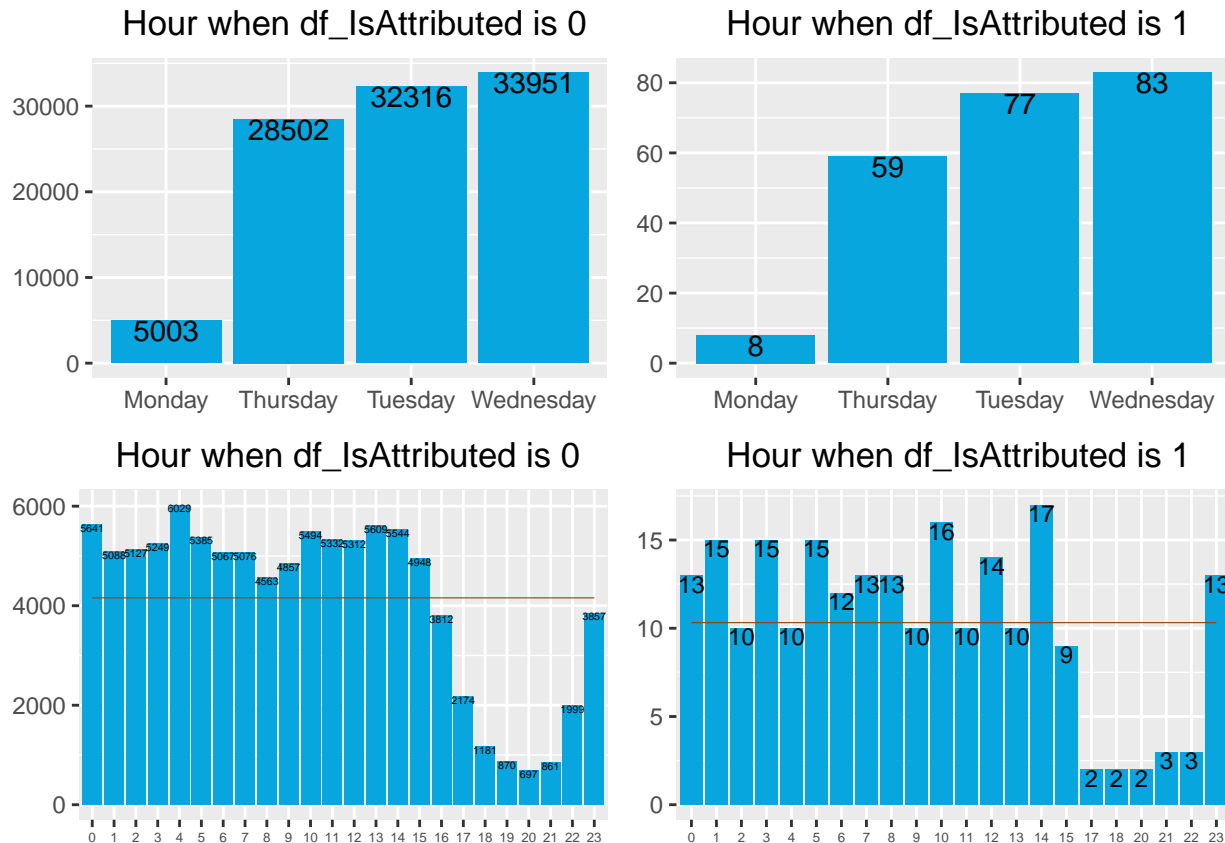
Hour when df_IsAttributed is 0



Hour when df_IsAttributed is 1



```
grid.arrange(p1_IsA_Day, p2_IsA_Day, p1_IsA_Hour, p2_IsA_Hour, ncol = 2)
```



There is a significative difference between that hour download was made.

It's interesting to use statistics to know more about this difference, but basically, the no-downloads happened more between 0 to 15h.

In this sample, there isn't download at the 16h, but there is a lot no-download.

2.2.2 investigating more the relations between day and time variables

is_attributed0 hour x day

```
df_IsAttributed0HourDay <- df_IsAttributed0 %>%
  group_by(weekday, hour) %>%
  summarise(counts = n())
```

`summarise()` has grouped output by 'weekday'. You can override using the `.groups` argument.

```
a <- ggplot(df_IsAttributed0HourDay, aes(x = hour, y = counts)) +
  geom_bar(fill = "#00A4DEF7", stat = "identity") +
  facet_grid(weekday ~ .) +
  geom_text(aes(label = counts), vjust = 1, size = 1.5) +
  ggtitle("is_attributed0 - Hour x weekday") +
  theme(axis.text.x = element_text(size=5),
        plot.title = element_text(hjust = 0.5))
```

is_attributed1 hour x day

```
df_IsAttributed1HourDay <- df_IsAttributed1 %>%
```

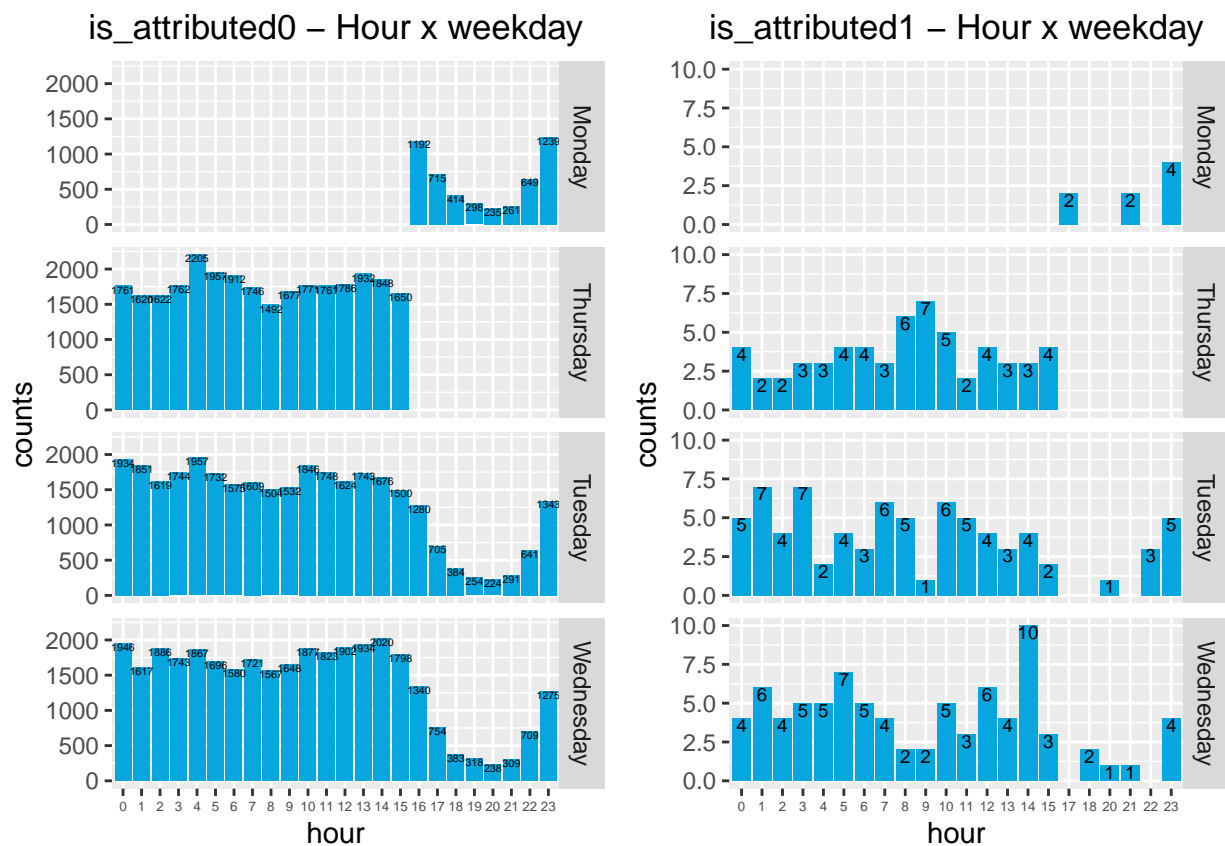


```
group_by(weekday, hour) %>%
summarise(counts = n())
```

`summarise()` has grouped output by 'weekday'. You can override using the `.groups` argument.

```
b <- ggplot(df_IsAttributed1HourDay, aes(x = hour, y = counts)) +
  geom_bar(fill = "#00A4DEF7", stat = "identity") +
  facet_grid(weekday ~ .) +
  geom_text(aes(label = counts), vjust = 1, size = 2.5) +
  ggtitle("is_attributed1 - Hour x weekday") +
  theme(axis.text.x = element_text(size=5),
        plot.title = element_text(hjust = 0.5))
```

```
# there wasn't downloads at Thursday in 17 to 03h
grid.arrange(a, b, ncol = 2)
```



Stage 3: Correlation

```
# Generating new columns from 'IP' relating it with another columns
data <- df %>%
  add_count(ip, weekday, hour) %>% rename("ipDayHour" = n) %>%
  add_count(ip, hour, channel) %>% rename("ipHourChannel" = n) %>%
  add_count(ip, hour, os) %>% rename("ipHourOs" = n) %>%
  add_count(ip, hour, app) %>% rename("ipHourApp" = n) %>%
  add_count(ip, hour, device) %>% rename("ipHourDevice" = n) %>%
```

```


select(-click_time)
# 'click_time' isn't necessary as we already have created weekday and hour columns

## Creating a dataset for doesn't alter the original
data_num <- data
# to observe how much each variable correlates depending if it was made download or no
data_0 <- data %>%
  filter(is_attributed == '0') %>%
  select(-is_attributed)
data_1 <- data %>%
  filter(is_attributed == '1') %>%
  select(-is_attributed)

# converting to numeric to correlate variables
data_num[,1:ncol(data_num)] = lapply(data_num[,1:ncol(data_num)], as.numeric)
data_0[,1:ncol(data_0)] <- lapply(data_0[,1:ncol(data_0)], as.numeric)
data_1[,1:ncol(data_1)] <- lapply(data_1[,1:ncol(data_1)], as.numeric)

## correlations
corrplot(cor(data_num, method="pearson"), method = 'number',
          number.cex= 7/ncol(data_num), type="lower")
mtext("data_num", at=12, line=2, cex=1)


```




```

corrplot(cor(data_0, method="pearson"), method = 'number',
          number.cex= 7/ncol(data_0), type="lower")
mtext("data_0", at=12, line=2, cex=1)

```



```
corrplot(cor(data_1, method="pearson"), method = 'number',
         number.cex= 7/ncol(data_1), type="lower")
mtext("data_1", at=12, line=2, cex=1)
```



data_num: only 'ip' and 'app' are very weakly and positively correlated with the target variable; 'device' and 'os' correlate positively with 'app' and with each other.

data_0: particularly, for no-downloads 'OS' is high positive correlated with 'device', and weakly is 'app' with them.

data_1: there are some weakly and very weakly correlations; 'ip', 'app', and 'os' aren't correlated as 'data_num' and 'data_0'